



# Enterprise Database Migration

Testing and Monitoring Databases  
in Google Cloud

Julianne Cuneo  
Data Analytics Specialist,  
Google Cloud



Hi, this is Julianne again. Welcome back to our course on Enterprise Database Migration.

Testing is important to ensure your migration is done correctly. Monitoring will help keep your migrated applications and databases running smoothly over time.

Now that you have your SQL Server instance in the cloud, let's take a look at some things you can do to monitor and test it.

## Learning objectives

- Use unit, integration, and regression testing techniques to ensure database migration success.
- Monitor your migration projects with Google tools.



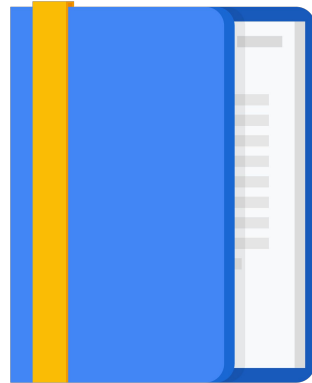
Taking a software development or DevOps approach to migration using unit, integration, and regression testing can ensure a successful migration process.

Aside from testing, There's also a need to -monitor- the entire migration, and Google has tools for that as well.

# Agenda

Testing

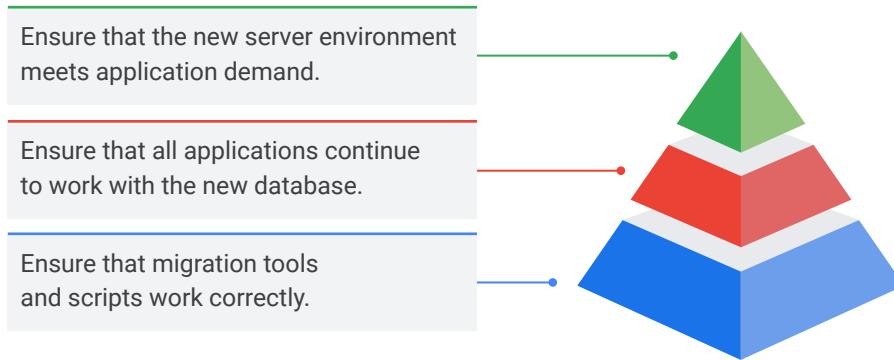
Monitoring



Thorough, automated testing is an important part of any IT project. Without effective testing there is no way to know if your database has been migrated correctly. There are many things you need to verify. Was the database schema migrated correctly? Has all the data been migrated? How about user logins? Can all the users still connect and can users only access the data they are permitted to access?

Let's start off by exploring testing.

## Thorough, automated testing can help uncover errors quickly and make migrations go faster



Complete testing is the surest way to facilitate a successful migration and can even speed up the process.

It is critical to ensure that applications continue to work on the new server and that the new server has sufficient resources to meet the demands of the application.

Before starting a long migration process, you want to make sure that the tools and scripts work properly.

---

Testing database migrations can be grouped into three broad categories

1

Structural  
testing

2

Functional  
testing

3

Non-functional  
testing

There are basically three categories of testing that need to be considered: structural, functional, and non-functional.

## Structural testing validates that the schema of the migrated database matches the source

### Test to ensure that:

- All tables have been migrated.
- Fields, data types, lengths, and constraints were migrated as expected.
- All indexes have been created.
- Primary-foreign key relationships were created.
- Primary and foreign keys match correctly in related tables.
- All stored procedures and triggers were migrated.
- User logins were migrated successfully.



Structural testing validates that everything that should have been moved has successfully been moved. It checks to make sure that every table, procedure, index, foreign key, trigger, login, and more has been brought over and that the new environment is structurally the same as the original.

Write automated structural tests to ensure that:

- All tables have been migrated.
- Fields, data types, lengths, and constraints were migrated as expected.
- All indexes have been created.
- Primary-foreign key relationships were created.
- Primary and foreign keys match correctly in related tables.
- All stored procedures and triggers were migrated.
- And lastly that User logins were migrated successfully.

---

## Functional testing ensures that all data was migrated and all database functions work as expected

### Test to ensure that:

- All data was migrated correctly.
- Stored procedures and triggers work as expected.
- User logins work as expected.
- Users can only perform operations and access data as permitted.
- Queries and views continue to return expected results.



Functional testing goes to the next level and makes sure that all of the data has come across correctly. It also makes sure that all of the views and stored procedures work as intended and that users can log in and see what they are supposed to, but not what they shouldn't.

Write automated functional tests to ensure that:

- All data was migrated correctly.
- Stored procedures and triggers work as expected.
- User logins work as expected.
- Users can only perform operations and access data as permitted.
- And that Queries and views continue to return expected results.

---

## Non-functional testing ensures that the new database system meets its business requirements

### **Types of non-functional tests include:**

- Performance testing to ensure that queries and transactions complete quickly enough to meet user expectations
- Load testing to ensure that the application performs at peak demand
- Stress testing to find the load at which the database breaks



After you know that everything works as intended, the next step is to make sure it performs well and can withstand peak usage periods. Load and stress testing will push the server to its limit to see whether it can handle the intended workload and just how far it will go before it breaks.

This helps determine whether more resources need to be allocated to the server.

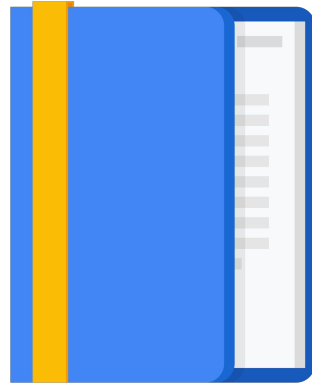


---

# Agenda

Testing

Monitoring



After a functional server is up and running, you want to monitor its activity to make sure there are no surprises.

## Cloud Operations monitoring collects and displays information about your resources

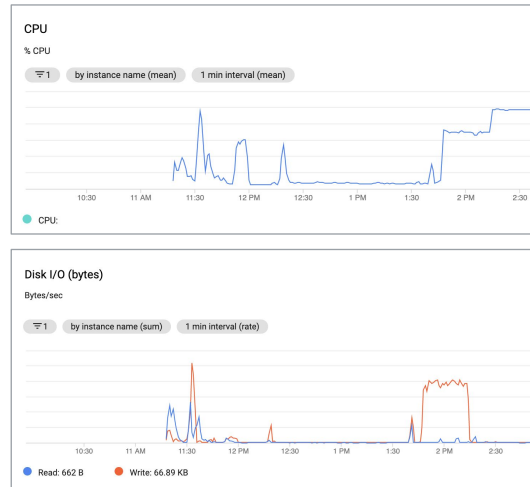
- Default monitoring is built into the web-based console for every service.
- You can create additional dashboards with custom charts.
- You can monitor resources in both AWS and Google Cloud.



Google provides a web-based monitoring tool with a default console for each service. You can create your own custom dashboards and monitor resources in both Google Cloud and AWS.

## Compute Engine VMs offer basic monitoring by default

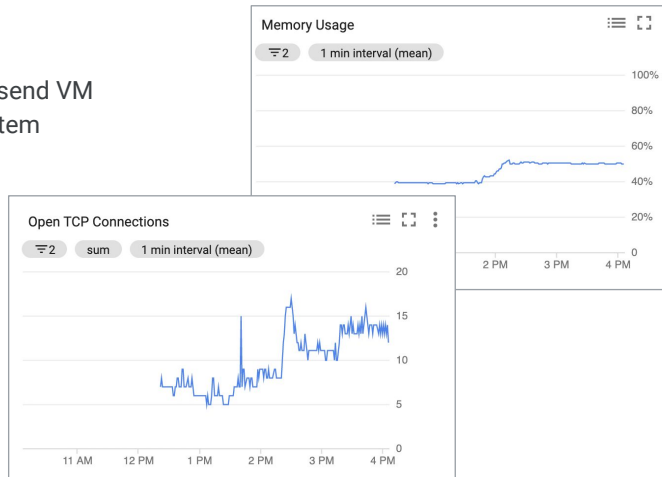
- Metrics include:
  - CPU usage
  - Network traffic
  - Disk I/O



There is basic VM monitoring by default, but that just shows “big-picture” stats such as CPU usage, network traffic, and Disk I/O.

## Install the Monitoring Agent to monitor additional metrics in a VM

- Uses a `collectd` daemon to send VM metrics to the Monitoring system
- Additional metrics include:
  - Memory usage
  - Swap usage
  - TCP connections



To really understand what's going on inside of a database server, you typically need to look at more metrics such as memory usage and swap file usage. You can install the Monitoring Agent to collect those metrics from the VM and send them to the Monitoring system.

The Monitoring Agent uses a “`collectd`” daemon to send VM metrics to the Monitoring system.

## The Monitoring Agent also monitors applications on virtual machines

- Many applications are supported on Linux, including:
    - Apache Web server
    - MySQL
    - Nginx
    - JVM
  - On Windows, the monitoring agent supports IIS and SQL Server
- SQL Server metrics include:
    - User connections
    - Transactions
    - Write transactions



The Monitoring Agent is flexible enough to monitor a variety of applications such as MySQL, Apache web server, and JVM. When you are running a Windows VM, it can monitor IIS and SQL Server.

This allows you to tap into the SQL Server-specific metrics usually provided with the SQL Server Performance Monitor, such as number of SQL connections, transactions, and cache hit ratio.

---

## On Windows, download and install the Monitoring Agent using Powershell

```
cd $env:UserProfile;  
(New-Object  
Net.WebClient).DownloadFile("https://repo.stackdriver.com/windows/StackdriverMonitoring-GCM-46.exe", ".\StackdriverMonitoring-GCM-46.exe")  
.\StackdriverMonitoring-GCM-46.exe
```



This Powershell command downloads and installs the Monitoring Agent on Windows. You can run it after logging on to the machine, or better still, make it part of the machine's startup script.

For more information, search Google Cloud documentation for the Monitoring agent.

From the documentation you can copy and paste the script shown.

## On Linux, download and install the Monitoring Agent using Bash

```
curl -sS0 https://dl.google.com/cloudagents/add-monitoring-agent-repo.sh  
sudo bash add-monitoring-agent-repo.sh  
sudo apt-get update  
sudo apt-get install stackdriver-agent  
sudo service stackdriver-agent start
```

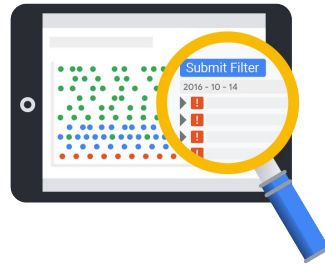


If you're running SQL Server on Linux instead, the bash commands shown here will install the monitoring agent. Make these commands part of your machine startup script.

Again, see the documentation to copy-paste the code.

## Logging allows you to store, search, analyze, monitor, and alert on log data

- Log data is collected and stored for each service.
- Logs can be filtered for analysis
- Log data can be exported to Cloud Storage, BigQuery, or Pub/Sub.
- Log metrics can be displayed in custom charts.
- You can create alerts from log metrics.



The agents collect the data from the OS and the SQL Server and store it in the Google Cloud logging service. Once there, it can be searched, filtered, and analyzed. You can export the log data to Cloud Storage, BigQuery, or Pub/Sub and even create alerts if the metrics meet a certain condition. Of course, you can also display custom charts of the metrics to visualize the data. You can even create log metrics and trigger alerts when there are anomalies in the log stream.



## Install the Logging Agent to monitor additional metrics in a VM

- Based on `fluentd`
- Streams logs from third-party applications and system software
- Linux and Windows

Query results			<a href="#">Jump to Now</a>	Actions ▾	Configure ▾
SEVERITY	TIMESTAMP	EDIT ▾	SUMMARY		
> ⓘ	2020-05-18 17:44:05.239 EDT		{ "reporter": "src", "start_time": "2020-05-18T17:44:05.239 EDT", "end_time": "2020-05-18T17:44:05.239 EDT", "src_instance": { ... }, "src_location": { ... }, "dest_instance": { ... }, "dest_location": { ... }, "bytes_sent": "0", "bytes_received": "0", "start_time": "2020-05-18T17:44:05.239 EDT", "end_time": "2020-05-18T17:44:05.239 EDT", "src_vpc": { ... }, "src_location": { ... }, "dest_vpc": { ... }, "dest_location": { ... }, "bytes_sent": "0", "bytes_received": "0" }		
> ⓘ	2020-05-18 17:44:05.239 EDT		{ "src_instance": { ... }, "bytes_sent": "0", "dest_instance": { ... }, "bytes_received": "0" }		
> ⓘ	2020-05-18 17:44:05.239 EDT		{ "src_location": { ... }, "dest_location": { ... }, "bytes_sent": "0", "bytes_received": "0" }		
> ⓘ	2020-05-18 17:44:05.239 EDT		{ "src_instance": { ... }, "bytes_sent": "168", "dest_instance": { ... }, "bytes_received": "0" }		
> ⓘ	2020-05-18 17:44:05.239 EDT		{ "start_time": "2020-05-18T17:44:03.202654568Z", "end_time": "2020-05-18T17:44:05.239 EDT", "src_instance": { ... }, "src_location": { ... }, "dest_instance": { ... }, "dest_location": { ... }, "bytes_sent": "168", "bytes_received": "0" }		
> ⓘ	2020-05-18 17:44:10.768 EDT		{ "src_vpc": { ... }, "src_location": { ... }, "dest_vpc": { ... }, "dest_location": { ... }, "bytes_sent": "0", "bytes_received": "0" }		
> ⓘ	2020-05-18 17:44:10.768 EDT		{ "src_instance": { ... }, "bytes_sent": "0", "dest_instance": { ... }, "bytes_received": "0" }		



You should install the Logging Agent if you are running your database on a virtual machine. The Logging Agent is based on fluentd and captures machine logs and the logs of third-party software, including SQL Server.

This works on both Windows and Linux.

---

## On Windows, install the Logging Agent using Powershell

```
cd $env:UserProfile;  
(New-Object  
Net.WebClient).DownloadFile("https://dl.google.com/cloudagents/windows/StackdriverLogging-v1-11.exe", ".\StackdriverLogging-v1-11.exe")  
.\StackdriverLogging-v1-11.exe
```



Here's the Powershell script to install the Logging Agent on Windows.

For more information, search Google Cloud documentation for the Logging agent.

## On Linux, install the Logging Agent using Bash

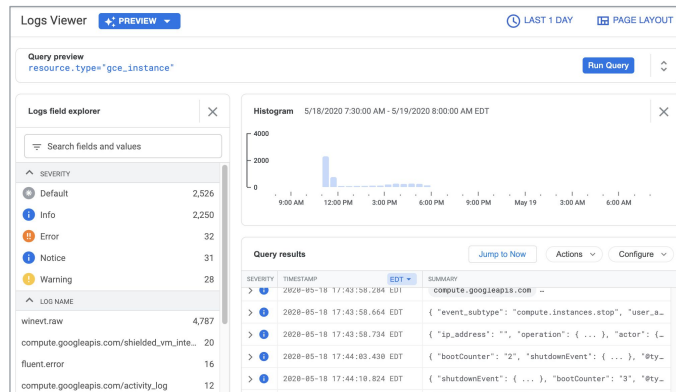
```
curl -sS0 https://dl.google.com/cloudagents/add-logging-agent-repo.sh  
sudo bash add-logging-agent-repo.sh  
sudo apt-get update  
sudo apt-get install google-fluentd  
sudo service google-fluentd start
```



Of course, you can also install the logging agent on Linux. Shown here are the equivalent bash commands. As before, copy and paste the code from the documentation.

## Use the Logs Viewer to filter and analyze the logs

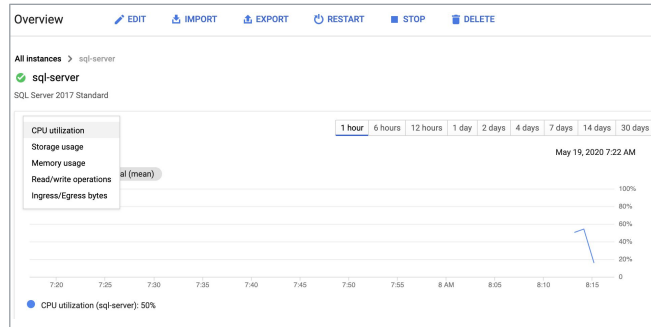
- Filter logs by service, resource, date/time range, severity, etc.
- Export logs to external sink (Cloud Storage, BigQuery, Pub/Sub).



The Logs Viewer provides a convenient web-based interface to filter the logged data, or you can export it and explore it using any other tools. To export logs, create a filter that collects the log information you are interested in, and then create an export sink. The export destination can be Cloud Storage, if you just want to save the log data, or BigQuery, if you want to use SQL for log analysis; or send the data to Pub/Sub for real-time log analysis.

## Cloud SQL provides monitoring by default

- No need to install the agents
- Metrics include:
  - CPU utilization
  - Storage usage
  - Memory usage
  - Read/write operations
  - Ingress/egress bytes
- Available from the Console and Monitor Dashboards



If you use Cloud SQL instead of building your own VMs, you get monitoring by default.

There's no need to install any agents. You get multiple standard metrics and can view this from the Console and the Monitor Dashboards.

The built-in metrics include:

CPU utilization  
Storage usage  
Memory usage  
Read/write operations  
and both ingress and egress bytes

These are the metrics you would probably want to monitor when managing a database.

## The Monitoring system automatically creates dashboards for resources in your projects

The screenshot shows the 'Dashboards Overview' page in Google Cloud Monitoring. At the top, there's a header with 'Dashboards Overview', a '+ CREATE DASHBOARD' button, and a 'BIGTABLE KEY VISUALIZER' link. Below the header, there's a 'Recent Dashboards' section with two cards: 'Microsoft SQL Server' (Application) and 'VM Instances' (Google Cloud Platform). Underneath, there's a 'Filter Dashboards' section with a search icon and a table of dashboards. The table has columns for 'Name' and 'Type'. The table lists six dashboards: Cloud SQL, Cloud Storage, Disks, Firewalls, Microsoft SQL Server, and VM Instances, all with star icons in the 'Name' column. The 'Type' column lists the associated platform for each resource.

Name	Type
Cloud SQL	Google Cloud Platform
Cloud Storage	Google Cloud Platform
Disks	Google Cloud Platform
Firewalls	Google Cloud Platform
Microsoft SQL Server	Application
VM Instances	Google Cloud Platform



The Monitoring system automatically creates dashboards for resources in your projects. It analyzes what resources have been created in a project and builds dashboards accordingly. This is a great way to get started. Often, the hardest part of monitoring is figuring out what you want to monitor.

After you get some experience, you can take this a step further by creating your own custom dashboards.

## Module review



In this module, you learned to use automated unit, integration, and regression testing techniques to ensure database migration success, and monitor your migration projects with Google tools.

